

Logestic_regression_and_synthetic_data

September 30, 2020

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import math

titanic=pd.read_csv(r"D:\msc3\machine learning\lab6\titanictrain.csv")
titanic.head(10)
```

```
[3]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3
5             6         0         3
6             7         0         1
7             8         0         3
8             9         1         3
9            10         1         2
```

```

                                Name      Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris    male  22.0     1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0     1
2                        Heikkinen, Miss. Laina    female  26.0     0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0     1
4                        Allen, Mr. William Henry    male  35.0     0
5                        Moran, Mr. James          male   NaN     0
6                        McCarthy, Mr. Timothy J    male  54.0     0
7                        Palsson, Master. Gosta Leonard    male   2.0     3
8  Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)  female  27.0     0
9                        Nasser, Mrs. Nicholas (Adele Achem)  female  14.0     1
```

```

Parch      Ticket      Fare Cabin Embarked
0      0    A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0  STON/O2. 3101282   7.9250   NaN        S
```

3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
5	0	330877	8.4583	NaN	Q
6	0	17463	51.8625	E46	S
7	1	349909	21.0750	NaN	S
8	2	347742	11.1333	NaN	S
9	0	237736	30.0708	NaN	C


```
[4]: print("Number of passangers:"+str(len(titanic.index)))
```

Number of passangers:891

0.1 Analyzing data


```
[5]: sns.countplot(x="Survived", data=titanic)
```

```
[5]: <matplotlib.axes._subplots.AxesSubplot at 0x194cb065188>
```




```
[6]: sns.countplot(x="Survived", hue="Sex", data=titanic)
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x194d64839c8>
```




```
[7]: sns.countplot(x="Survived", hue="Pclass", data=titanic)
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x194d65102c8>
```




```
[8]: titanic["Age"].plot.hist()
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x194d65916c8>
```



```
[9]: titanic["Fare"].plot.hist(bins=20,figsize=(10,5))
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x194d6626488>
```




```
[10]: titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   PassengerId     891 non-null   int64  
1   Survived        891 non-null   int64  
2   Pclass          891 non-null   int64  
3   Name            891 non-null   object  
4   Sex             891 non-null   object  
5   Age             714 non-null   float64  
6   SibSp           891 non-null   int64  
7   Parch           891 non-null   int64  
8   Ticket          891 non-null   object  
9   Fare            891 non-null   float64  
10  Cabin           204 non-null   object  
11  Embarked        889 non-null   object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

```
[11]: sns.countplot(x="SibSp", data=titanic)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x194d6711288>
```



0.2 Data Wrangling

```
[12]: titanic.isnull()
```

```
[12]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	\
0	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	
..	
886	False	False	False	False	False	False	False	False	False	
887	False	False	False	False	False	False	False	False	False	
888	False	False	False	False	False	True	False	False	False	
889	False	False	False	False	False	False	False	False	False	
890	False	False	False	False	False	False	False	False	False	

	Fare	Cabin	Embarked
0	False	True	False
1	False	False	False
2	False	True	False
3	False	False	False
4	False	True	False
..
886	False	True	False
887	False	False	False
888	False	True	False
889	False	False	False
890	False	True	False

[891 rows x 12 columns]


```
[13]: titanic.isnull().sum()
```

```
[13]: PassengerId      0
Survived             0
Pclass               0
Name                 0
Sex                  0
Age                 177
SibSp                0
Parch                0
Ticket               0
Fare                 0
```

```
Cabin          687
Embarked       2
dtype: int64
```

```
[14]: sns.boxplot(x="Pclass", y="Age", data=titanic)
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x194d6780108>
```



```
[15]: titanic.head(5)
```

```
[15]:   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
```

```
      Name      Sex  Age  SibSp  \
0  Braund, Mr. Owen Harris    male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0    1
2    Heikkinen, Miss. Laina  female  26.0    0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0    1
4    Allen, Mr. William Henry    male  35.0    0
```

```
   Parch      Ticket    Fare Cabin Embarked
```

0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[16]: titanic.drop("Cabin", axis=1, inplace=True)
```

```
[17]: titanic.head()
```

```
[17]: PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Embarked
0	0	A/5 21171	7.2500	S
1	0	PC 17599	71.2833	C
2	0	STON/O2. 3101282	7.9250	S
3	0	113803	53.1000	S
4	0	373450	8.0500	S

```
[18]: titanic.dropna(inplace=True)
```

```
[19]: titanic.isnull().sum()
```

```
[19]: PassengerId    0
Survived          0
Pclass            0
Name              0
Sex               0
Age               0
SibSp             0
Parch             0
Ticket            0
Fare              0
Embarked          0
dtype: int64
```



```
[20]: sex=pd.get_dummies(titanic['Sex'], drop_first=True)
sex.head()
```

```
[20]:      male
0      1
1      0
2      0
3      0
4      1
```

```
[21]: embark=pd.get_dummies(titanic['Embarked'], drop_first=True)
embark.head()
```

```
[21]:      Q  S
0  0  1
1  0  0
2  0  1
3  0  1
4  0  1
```

```
[22]: titanic=pd.concat((titanic,sex,embark),axis=1)
```

```
[23]: titanic.head()
```

```
[23]:      PassengerId  Survived  Pclass  \
0              1         0         3
1              2         1         1
2              3         1         3
3              4         1         1
4              5         0         3
```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                Heikkinen, Miss. Laina   female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                Allen, Mr. William Henry   male  35.0      0
```

```

      Parch      Ticket    Fare Embarked  male  Q  S
0      0      A/5 21171    7.2500      S    1  0  1
1      0      PC 17599   71.2833      C    0  0  0
2      0  STON/O2. 3101282    7.9250      S    0  0  1
3      0      113803   53.1000      S    0  0  1
4      0      373450    8.0500      S    1  0  1
```

```
[24]: titanic.
      ↳drop(['Sex', 'Name', 'PassengerId', 'Ticket', 'Embarked'],axis=1,inplace=True)
```

```
[25]: titanic.head()
```

```
[25]:   Survived  Pclass   Age  SibSp  Parch    Fare   male  Q  S
0         0       3  22.0     1     0   7.2500     1  0  1
1         1       1  38.0     1     0  71.2833     0  0  0
2         1       3  26.0     0     0   7.9250     0  0  1
3         1       1  35.0     1     0  53.1000     0  0  1
4         0       3  35.0     0     0   8.0500     1  0  1
```

0.3 Train and Test

```
[26]: X=titanic.drop("Survived",axis=1)
      y=titanic["Survived"]
```

```
[27]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,
      ↪random_state=2)
```

```
[28]: from sklearn.linear_model import LogisticRegression
      logreg = LogisticRegression(C=1e5)
      logreg.fit(X_train,y_train)
      predictions=logreg.predict(X_test)
```

C:\Users\blr0a\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:764: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
[29]: from sklearn.metrics import classification_report
      classification_report(y_test,predictions)
```

```
[29]: '          precision    recall  f1-score   support\n\n 0.82      0.79      0.81      130\n 84\n\n 0.76      0.77      0.76      214\n214\n\n          accuracy          0.77      214\nmacro avg          0.77      0.77      0.77      214\nweighted avg          0.77      0.77      0.77      214\n'
```

```
[30]: from sklearn.metrics import confusion_matrix
      confusion_matrix(y_test,predictions)
```

```
[30]: array([[103, 27],
           [ 22, 62]], dtype=int64)
```


```
[31]: from sklearn.metrics import accuracy_score
print("train accuracy:")
print(format(logreg.score(X_train,y_train)*100.0))
print("test accuracy:")
accuracy_score(y_test,predictions)
```

train accuracy:
80.72289156626506
test accuracy:

```
[31]: 0.7710280373831776
```

0.4 USING make_classification

```
[32]: from sklearn.datasets import make_classification as mc
# Generate Clean data
X,y = mc(n_samples=10000, n_features=2, n_informative=2,n_redundant=0,
        ↪n_repeated=0, n_classes=2,
        ↪n_clusters_per_class=1,class_sep=2,flip_y=0,weights=[0.5,0.5],
        ↪random_state=17)
f, (ax1,ax2) = plt.subplots(nrows=1, ncols=2,figsize=(20,8))
sns.scatterplot(X[:,0],X[:,1],hue=y,ax=ax1);
ax1.set_title("No Noise");
# Generate noisy Data
X,y = mc(n_samples=10000, n_features=2, n_informative=2, n_redundant=0,
        ↪n_repeated=0, n_classes=2, n_clusters_per_class=1,class_sep=2,flip_y=0.
        ↪2,weights=[0.5,0.5], random_state=17)
sns.scatterplot(X[:,0],X[:,1],hue=y,ax=ax2);
ax2.set_title("With Noise");
plt.show();
```



0.5 applying logistic regression on make_classification(Noiseless data)

```
[33]: x,y=mc(n_samples=10000, n_features=8,n_redundant=0, n_repeated=0, n_classes=3,  
↳n_clusters_per_class=1, random_state=17)
```

```
[100]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA  
lda1=LDA(n_components=2)  
x_lda=lda1.fit_transform(x,y)
```

```
[101]: from sklearn.linear_model import LogisticRegression as LR  
lr=LR()  
x_train,x_test,y_train,y_test=train_test_split(x_lda,y,random_state=0)  
l=lr.fit(x_train,y_train)  
coef=l.coef_  
intercept=l.intercept_
```

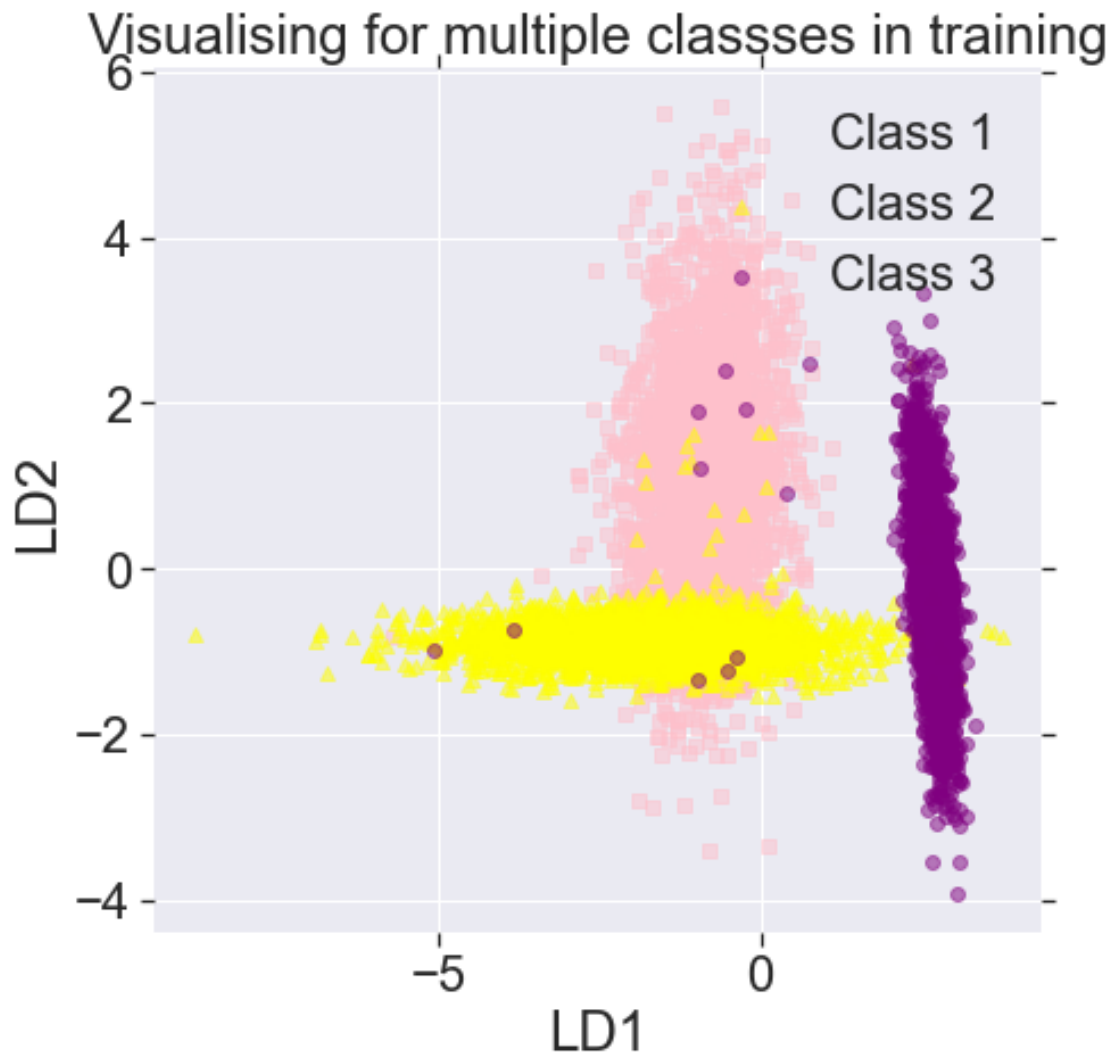
```
[102]: label_dict = {0:'Class 1', 1: 'Class 2', 2: 'Class 3', 3:'Class 4'}  
plt.figure(figsize=(7,7))  
  
def plot_scikit_lda(X,y, title):  
    ax = plt.subplot(111)  
    for label,marker,color in zip(  
        [i for i in range(4)],('s','^','o'),('pink','yellow','purple')):  
  
        plt.scatter(x=X[:,0][y == label],  
                    y=X[:,1][y == label],  
                    marker=marker,  
                    color=color,  
                    alpha=0.5,  
                    label=label_dict[label])  
  
    plt.xlabel('LD1')  
    plt.ylabel('LD2')  
  
    leg = plt.legend(loc='upper right', fancybox=True)  
    leg.get_frame().set_alpha(0)  
    plt.title(title)  
  
    # hide axis ticks  
    plt.tick_params(axis="both", which="both", bottom="off", top="off",  
                    labelbottom="on", left="off", right="off", labelleft="on")  
  
    # remove axis spines  
    ax.spines["top"].set_visible(False)
```

```

ax.spines["right"].set_visible(False)
ax.spines["bottom"].set_visible(False)
ax.spines["left"].set_visible(False)
plt.tight_layout
plt.show()

plot_scikit_lda(x_train, y_train, title='Visualising for multiple classes in_
→training')

```



```

[103]: y_pred=lr.predict(x_test)
        print(classification_report(y_true=y_test,y_pred=y_pred))

```

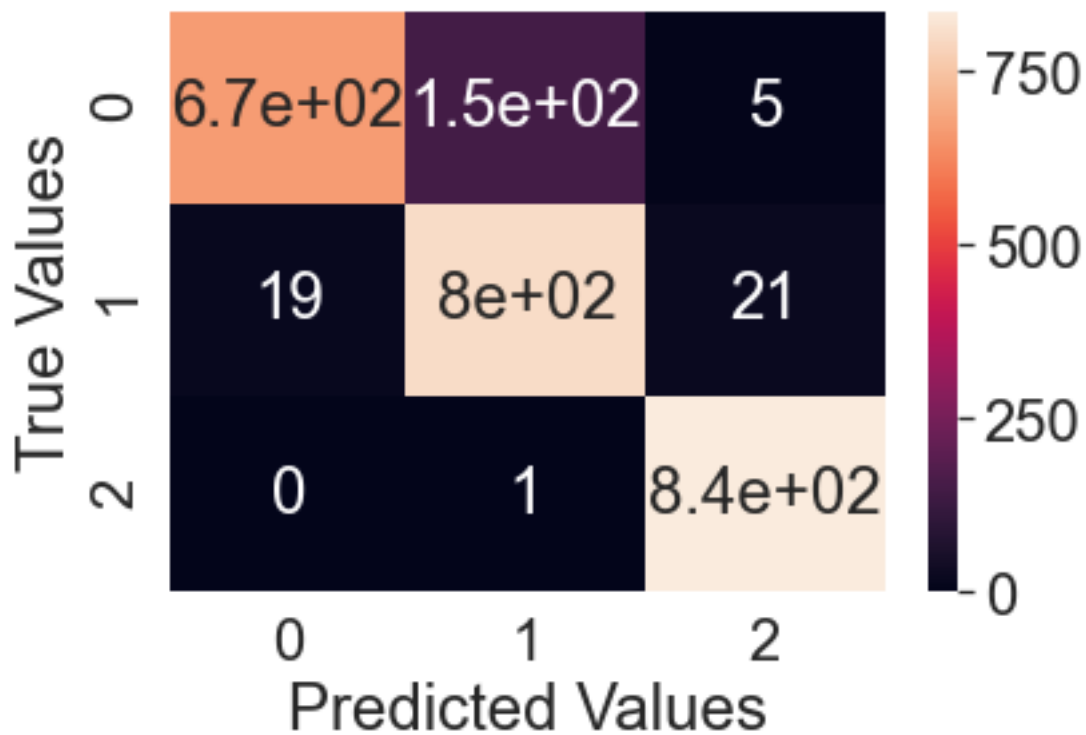
```

precision    recall  f1-score   support

```

0	0.97	0.81	0.88	822
1	0.84	0.95	0.89	841
2	0.97	1.00	0.98	837
accuracy			0.92	2500
macro avg	0.93	0.92	0.92	2500
weighted avg	0.93	0.92	0.92	2500

```
[104]: sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
plt.xlabel('Predicted Values')
plt.ylabel('True Values')
plt.show()
```



0.6 applying logistic regression on make_classification (noisy data)

```
[105]: x,y=mc(n_samples=10000, n_features=8,n_redundant=0, n_repeated=0, n_classes=3,
↳ n_clusters_per_class=1, random_state=17, weights=[0.3,0.3,0.4],flip_y=0.2)
```

```
[106]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
lda1=LDA(n_components=2)
x_lda=lda1.fit_transform(x,y)
```

```

[107]: lr=LR()
x_train,x_test,y_train,y_test=train_test_split(x_lda,y,random_state=0)
l=lr.fit(x_train,y_train)
coef=l.coef_
intercept=l.intercept_

[108]: label_dict = {0:'Class 1', 1: 'Class 2', 2: 'Class 3', 3:'Class 4'}
plt.figure(figsize=(7,7))

def plot_scikit_lda(X,y, title):
    ax = plt.subplot(111)
    for label,marker,color in zip(
        [i for i in range(4)],('*', '^', 'o'), ('black', 'yellow', 'purple')):

        plt.scatter(x=X[:,0][y == label],
                    y=X[:,1][y == label],
                    marker=marker,
                    color=color,
                    alpha=0.5,
                    label=label_dict[label])

    plt.xlabel('LD1')
    plt.ylabel('LD2')

    leg = plt.legend(loc='upper right', fancybox=True)
    leg.get_frame().set_alpha(0)
    plt.title(title)


    # hide axis ticks
    plt.tick_params(axis="both", which="both", bottom="off", top="off",
                    labelbottom="on", left="off", right="off", labelleft="on")

    # remove axis spines
    ax.spines["top"].set_visible(False)
    ax.spines["right"].set_visible(False)
    ax.spines["bottom"].set_visible(False)
    ax.spines["left"].set_visible(False)
    plt.tight_layout
    plt.show()

plot_scikit_lda(x_train, y_train, title='Visualising for multiple classes in_
↳training')

```

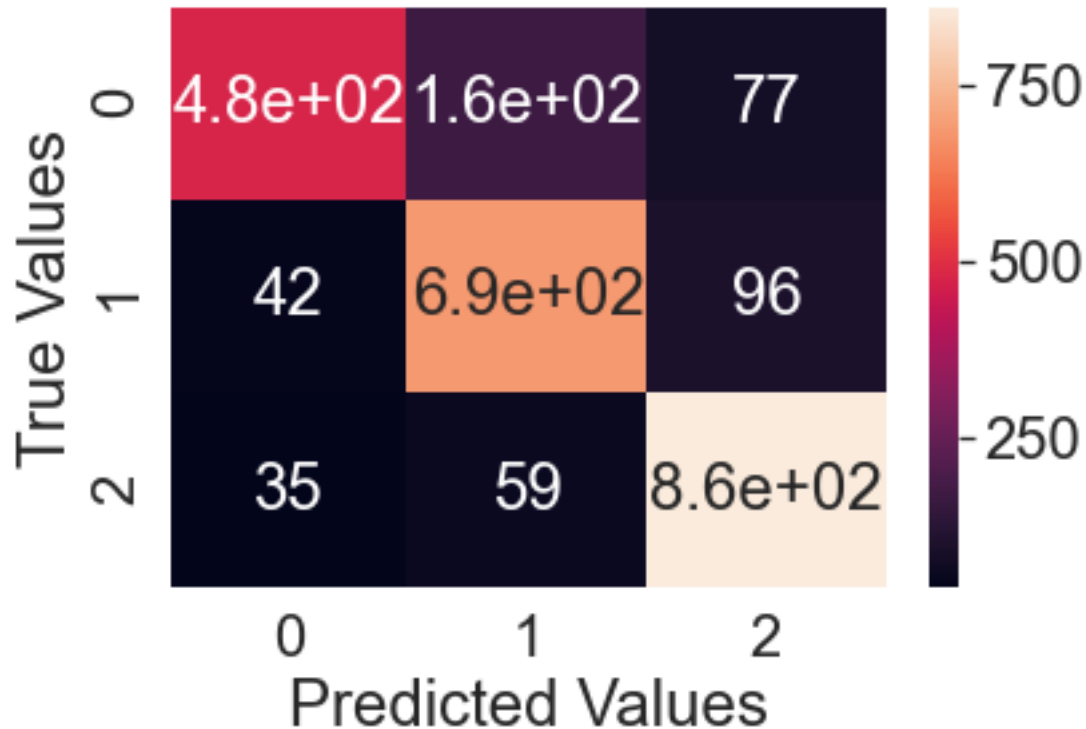
Visualising for multiple classes in training



```
[109]: y_pred=lr.predict(x_test)
print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.66	0.75	720
1	0.75	0.83	0.79	826
2	0.83	0.90	0.87	954
accuracy			0.81	2500
macro avg	0.82	0.80	0.80	2500
weighted avg	0.81	0.81	0.81	2500


```
[110]: sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
plt.xlabel('Predicted Values')
plt.ylabel('True Values')
plt.show()
```



0.7 applying logestic regression on make_multilabel_classification

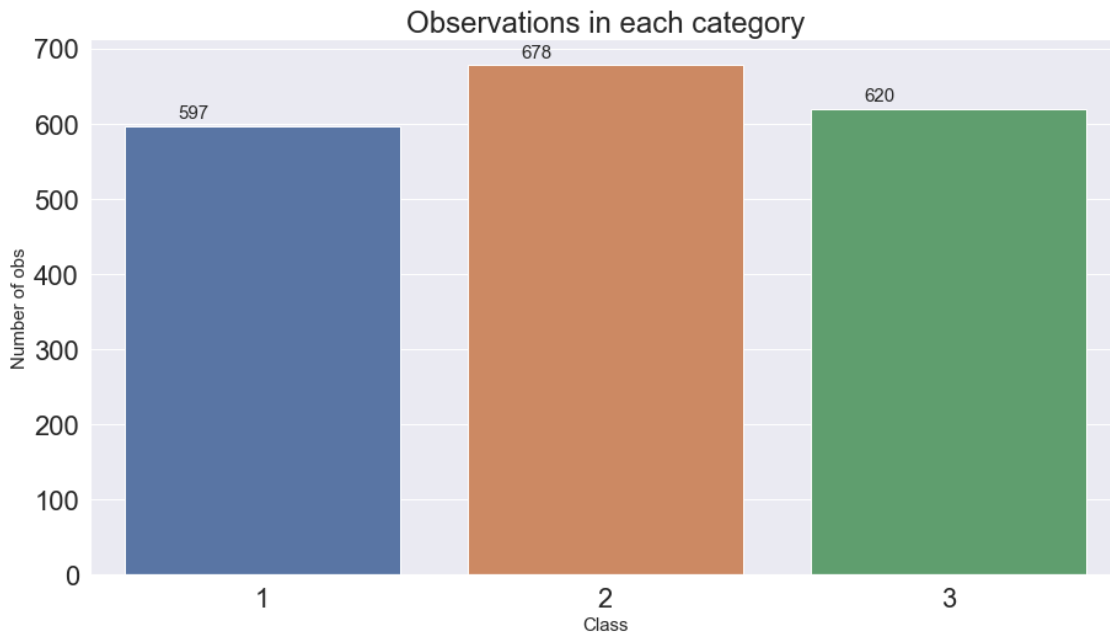
```
[111]: from sklearn.datasets import make_multilabel_classification as mmc
X,Y=mmc(n_samples=1000,n_features=8,n_classes=3,n_labels=2,length=30,random_state=0,allow_unlabeled=True)
```

```
[112]: categories = [i for i in range(1,4)]
sns.set(font_scale = 2)
plt.figure(figsize=(15,8))
ax= sns.barplot(categories, np.sum(Y[:,:],axis=0))
plt.title("Observations in each category", fontsize=24)
plt.ylabel('Number of obs', fontsize=15)
plt.xlabel('Class', fontsize=15)
#adding the text labels
rects = ax.patches
labels = np.sum(Y[:,:],axis=0)
for rect, label in zip(rects, labels):
    height = rect.get_height()
```

```

    ax.text(rect.get_x() + rect.get_width()/4, height + 5, label, ha='center',
    ↪va='bottom', fontsize=15)
plt.show()

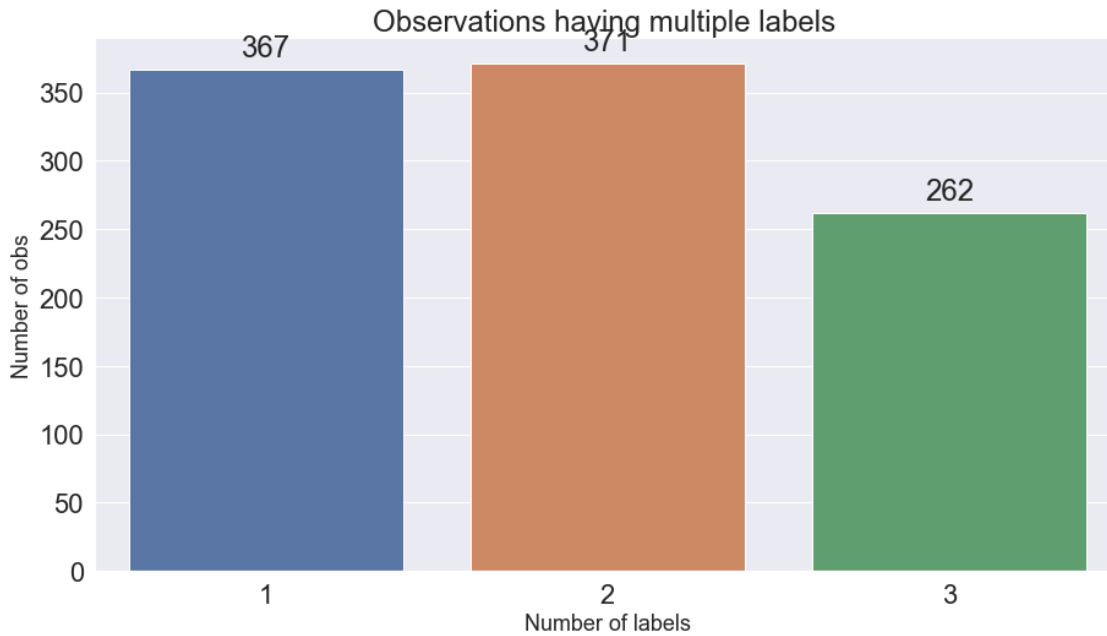
```



```

[113]: rowSums = np.sum(Y[:,:],axis=1)
multiLabel_counts = np.array([np.count_nonzero(rowSums==i) for i in range(1,4)])
sns.set(font_scale = 2)
plt.figure(figsize=(15,8))
ax = sns.barplot([i for i in range(1,4)], multiLabel_counts)
plt.title("Observations having multiple labels ")
plt.ylabel('Number of obs', fontsize=18)
plt.xlabel('Number of labels', fontsize=18)
#adding the text labels
rects = ax.patches
labels = multiLabel_counts
for rect, label in zip(rects, labels):
    height = rect.get_height()
    ax.text(rect.get_x() + rect.get_width()/2, height + 5, label, ha='center',
    ↪va='bottom')
plt.show()

```



```
[193]: label_dict = {0: 'Class 1', 1: 'Class 2', 2: 'Class 3'}
plt.figure(figsize=(15,15))

def plot_scikit_lda(X, title):
    for i in range(3):
        lda=LDA()
        X_lda=lda.fit_transform(X,Y[:,i])
        #print(X_lda.shape)
        ax = plt.subplot(2,2,i+1)
        for label,marker,color in zip(
            [0,1],('o','*','s'),('yellow','blue','red')):
            plt.scatter(x=X_lda[:,0][Y[:,i] == label],
                        y=np.ones(X_lda[Y[:,i] == label].shape[0]),#X[:,1][Y[:,i]
↳ == label],
                        marker=marker,
                        color=color,
                        alpha=0.5,
                        label=("rest",label_dict[i])[int(label==1)])

        plt.xlabel('LD1')
        plt.ylabel('LD2')

        leg = plt.legend(loc='upper right', fancybox=True)
        leg.get_frame().set_alpha(0)
        plt.title(title+label_dict[i])
```

```

# hide axis ticks
plt.tick_params(axis="both", which="both", bottom="off", top="off",
                labelbottom="on", left="off", right="off", labelleft="on")

# remove axis spines
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)
ax.spines["bottom"].set_visible(False)
ax.spines["left"].set_visible(False)
plt.tight_layout(pad=3.0)
plt.show()

plot_scikit_lda(X, title='Visualising for ')

```

```

↳ -----

ValueError                                Traceback (most recent call↳
↳ last)

<ipython-input-193-e569071f7855> in <module>
    37     plt.show()
    38
--> 39 plot_scikit_lda(X, title='Visualising for ')

<ipython-input-193-e569071f7855> in plot_scikit_lda(X, title)
     5     for i in range(3):
     6         lda=LDA()
----> 7         X_lda=lda.fit_transform(X,Y[:,i])
     8         #print(X_lda.shape)
     9         ax = plt.subplot(2,2,i+1)

~\Anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X,↳
↳ y, **fit_params)
    691         else:
    692             # fit method of arity 2 (supervised transformation)
--> 693             return self.fit(X, y, **fit_params).transform(X)
    694
    695

~\Anaconda3\lib\site-packages\sklearn\discriminant_analysis.py in↳
↳ fit(self, X, y)

```

```

423         """
424         X, y = self._validate_data(X, y, ensure_min_samples=2,
↪estimator=self,
--> 425                                     dtype=[np.float64, np.float32])
426         self.classes_ = unique_labels(y)
427         n_samples, _ = X.shape

~\Anaconda3\lib\site-packages\sklearn\base.py in _validate_data(self, X,
↪y, reset, validate_separately, **check_params)
430         y = check_array(y, **check_y_params)
431         else:
--> 432         X, y = check_X_y(X, y, **check_params)
433         out = X, y
434

~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
↪inner_f(*args, **kwargs)
71         FutureWarning)
72         kwargs.update({k: arg for k, arg in zip(sig.parameters,
↪args)})
---> 73         return f(**kwargs)
74         return inner_f
75

~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
↪check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy,
↪force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples,
↪ensure_min_features, y_numeric, estimator)
811         y = y.astype(np.float64)
812
--> 813         check_consistent_length(X, y)
814
815         return X, y

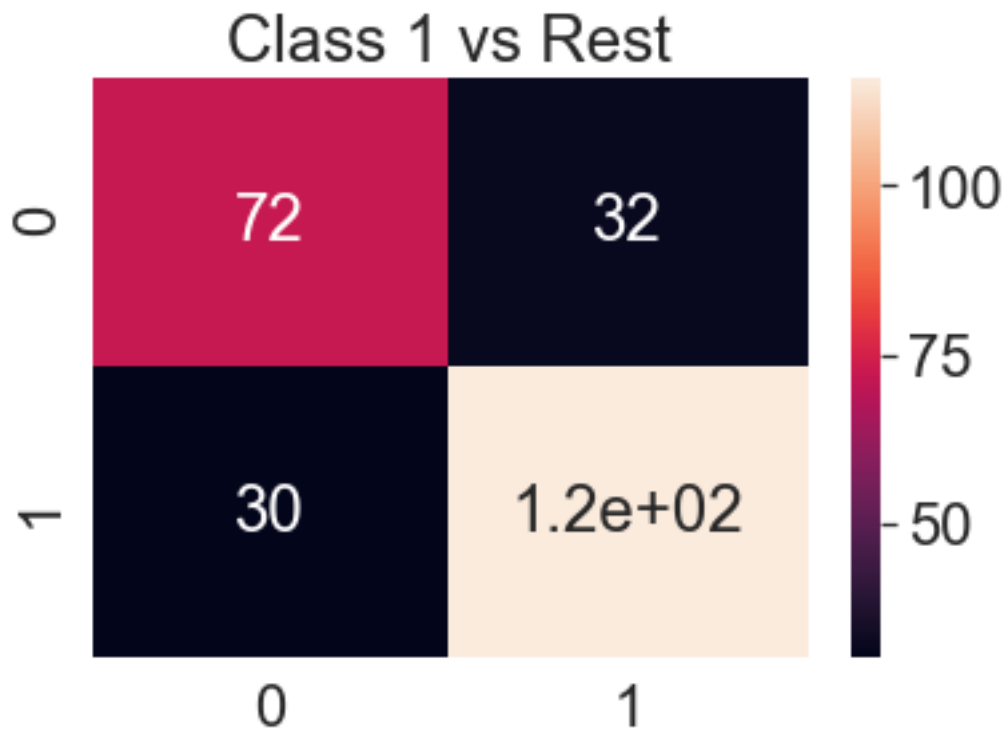
~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in
↪check_consistent_length(*arrays)
255         if len(uniques) > 1:
256             raise ValueError("Found input variables with inconsistent
↪numbers of"
--> 257                                     " samples: %r" % [int(l) for l in lengths])
258
259

```

ValueError: Found input variables with inconsistent numbers of samples:
→ [712, 1000]


<Figure size 1080x1080 with 0 Axes>

```
[115]: X_train,X_test,Y_train,Y_test=train_test_split(X,Y,random_state=0)
pred=np.ndarray(Y_test.shape)
for i in range(Y_test.shape[1]):
    lr.fit(X_train,Y_train[:,i])
    pred[:,i]=lr.predict(X_test)
    sns.heatmap(confusion_matrix(Y_test[:,i],pred[:,i]),annot=True)
    plt.title("Class "+str(i+1)+" vs Rest")
    plt.show()
    print(classification_report(Y_test[:,i],pred[:,i]))
```

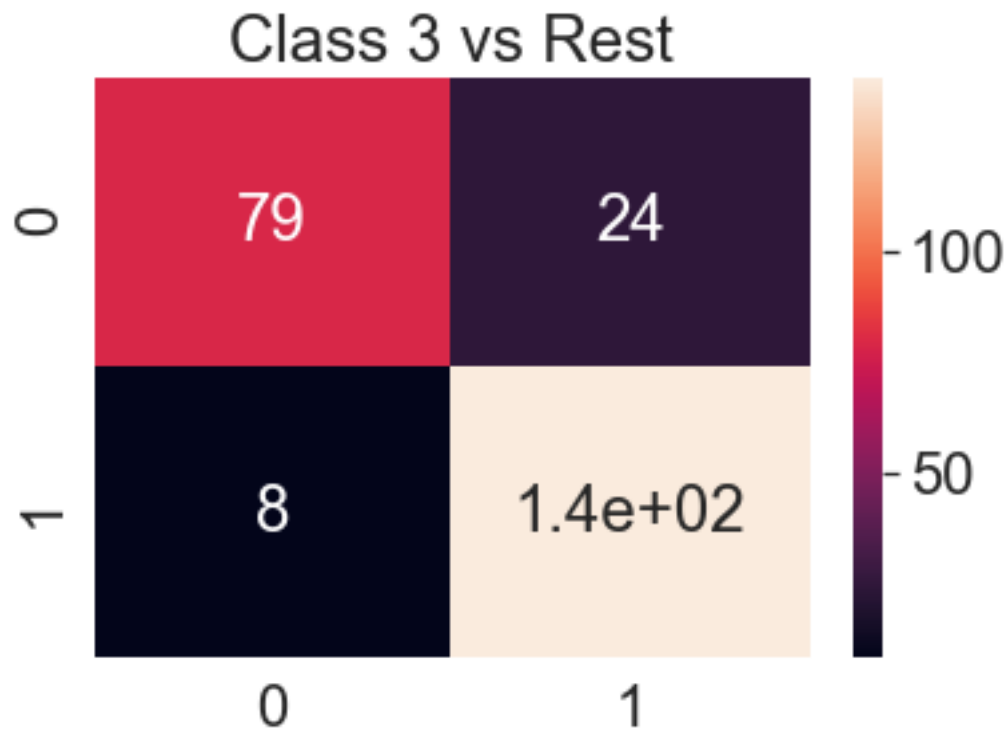


	precision	recall	f1-score	support
0	0.71	0.69	0.70	104
1	0.78	0.79	0.79	146
accuracy			0.75	250

macro avg	0.74	0.74	0.74	250
weighted avg	0.75	0.75	0.75	250



	precision	recall	f1-score	support
0	0.79	0.56	0.66	78
1	0.82	0.93	0.87	172
accuracy			0.82	250
macro avg	0.81	0.75	0.77	250
weighted avg	0.81	0.82	0.81	250



	precision	recall	f1-score	support
0	0.91	0.77	0.83	103
1	0.85	0.95	0.90	147
accuracy			0.87	250
macro avg	0.88	0.86	0.86	250
weighted avg	0.88	0.87	0.87	250

[]:

[]: