

Logistic_regression

September 30, 2020

1 Regularization with logestic regression

Regularization helps to solve over fitting problem in machine learning. Simple model will be a very poor generalization of data. At the same time, complex model may not perform well in test data due to over fitting. We need to choose the right model in between simple and complex model. Regularization helps to choose preferred model complexity, so that model is better at predicting. Regularization is nothing but adding a penalty term to the objective function and control the model complexity using that penalty term. It can be used for many machine learning algorithms.

Regularization of linear models Regularization is a method for “constraining” or “regularizing” the size of the coefficients, thus “shrinking” them towards zero. It reduces model variance and thus minimizes overfitting. If the model is too complex, it tends to reduce variance more than it increases bias, resulting in a model that is more likely to generalize.

Our aim is to locate the optimum model complexity, and thus regularization is useful when we believe our model is too complex.

1.1 Logestic regression without regularization

```
[30]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.metrics import confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
import seaborn as sns
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
```

```
[31]: breast_cancer=pd.read_csv(r"D:\msc3\machine learning\lab7\data.csv")
breast_cancer.head(10)
```

```
[31]:
```

| | id | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | \ |
|---|----------|-----------|-------------|--------------|----------------|-----------|---|
| 0 | 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 | |
| 1 | 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 | |
| 2 | 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 | |

| | | | | | | |
|---|----------|---|-------|-------|--------|--------|
| 3 | 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 |
| 4 | 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 |
| 5 | 843786 | M | 12.45 | 15.70 | 82.57 | 477.1 |
| 6 | 844359 | M | 18.25 | 19.98 | 119.60 | 1040.0 |
| 7 | 84458202 | M | 13.71 | 20.83 | 90.20 | 577.9 |
| 8 | 844981 | M | 13.00 | 21.82 | 87.50 | 519.8 |
| 9 | 84501001 | M | 12.46 | 24.04 | 83.97 | 475.9 |

| | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | \ |
|---|-----------------|------------------|----------------|---------------------|---|
| 0 | 0.11840 | 0.27760 | 0.30010 | 0.14710 | |
| 1 | 0.08474 | 0.07864 | 0.08690 | 0.07017 | |
| 2 | 0.10960 | 0.15990 | 0.19740 | 0.12790 | |
| 3 | 0.14250 | 0.28390 | 0.24140 | 0.10520 | |
| 4 | 0.10030 | 0.13280 | 0.19800 | 0.10430 | |
| 5 | 0.12780 | 0.17000 | 0.15780 | 0.08089 | |
| 6 | 0.09463 | 0.10900 | 0.11270 | 0.07400 | |
| 7 | 0.11890 | 0.16450 | 0.09366 | 0.05985 | |
| 8 | 0.12730 | 0.19320 | 0.18590 | 0.09353 | |
| 9 | 0.11860 | 0.23960 | 0.22730 | 0.08543 | |

| ... | radius_worst | texture_worst | perimeter_worst | area_worst | \ |
|-----|--------------|---------------|-----------------|------------|---|
| 0 | 25.38 | 17.33 | 184.60 | 2019.0 | |
| 1 | 24.99 | 23.41 | 158.80 | 1956.0 | |
| 2 | 23.57 | 25.53 | 152.50 | 1709.0 | |
| 3 | 14.91 | 26.50 | 98.87 | 567.7 | |
| 4 | 22.54 | 16.67 | 152.20 | 1575.0 | |
| 5 | 15.47 | 23.75 | 103.40 | 741.6 | |
| 6 | 22.88 | 27.66 | 153.20 | 1606.0 | |
| 7 | 17.06 | 28.14 | 110.60 | 897.0 | |
| 8 | 15.49 | 30.73 | 106.20 | 739.3 | |
| 9 | 15.09 | 40.68 | 97.65 | 711.4 | |

| | smoothness_worst | compactness_worst | concavity_worst | concave points_worst | \ |
|---|------------------|-------------------|-----------------|----------------------|---|
| 0 | 0.1622 | 0.6656 | 0.7119 | 0.2654 | |
| 1 | 0.1238 | 0.1866 | 0.2416 | 0.1860 | |
| 2 | 0.1444 | 0.4245 | 0.4504 | 0.2430 | |
| 3 | 0.2098 | 0.8663 | 0.6869 | 0.2575 | |
| 4 | 0.1374 | 0.2050 | 0.4000 | 0.1625 | |
| 5 | 0.1791 | 0.5249 | 0.5355 | 0.1741 | |
| 6 | 0.1442 | 0.2576 | 0.3784 | 0.1932 | |
| 7 | 0.1654 | 0.3682 | 0.2678 | 0.1556 | |
| 8 | 0.1703 | 0.5401 | 0.5390 | 0.2060 | |
| 9 | 0.1853 | 1.0580 | 1.1050 | 0.2210 | |

| | symmetry_worst | fractal_dimension_worst |
|---|----------------|-------------------------|
| 0 | 0.4601 | 0.11890 |
| 1 | 0.2750 | 0.08902 |

| | | |
|---|--------|---------|
| 2 | 0.3613 | 0.08758 |
| 3 | 0.6638 | 0.17300 |
| 4 | 0.2364 | 0.07678 |
| 5 | 0.3985 | 0.12440 |
| 6 | 0.3063 | 0.08368 |
| 7 | 0.3196 | 0.11510 |
| 8 | 0.4378 | 0.10720 |
| 9 | 0.4366 | 0.20750 |

[10 rows x 32 columns]

```
[32]: print("Number of data:"+str(len(breast_cancer.index)))
```

Number of data:569

```
[33]: breast_cancer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
```

```

26 smoothness_worst      569 non-null    float64
27 compactness_worst     569 non-null    float64
28 concavity_worst       569 non-null    float64
29 concave points_worst  569 non-null    float64
30 symmetry_worst        569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```

```
[34]: breast_cancer.isnull().any()
```

```


[34]: id                False
      diagnosis          False
      radius_mean        False
      texture_mean        False
      perimeter_mean      False
      area_mean           False
      smoothness_mean     False
      compactness_mean    False
      concavity_mean       False
      concave points_mean False
      symmetry_mean        False
      fractal_dimension_mean False
      radius_se           False
      texture_se          False
      perimeter_se        False
      area_se             False
      smoothness_se       False
      compactness_se      False
      concavity_se        False
      concave points_se   False
      symmetry_se         False
      fractal_dimension_se False
      radius_worst        False
      texture_worst       False
      perimeter_worst     False
      area_worst          False
      smoothness_worst    False
      compactness_worst   False
      concavity_worst     False
      concave points_worst False
      symmetry_worst       False
      fractal_dimension_worst False
      dtype: bool

```

2 Analyzing data

```
[35]: sns.countplot("diagnosis",data=breast_cancer)
```

```
[35]: <matplotlib.axes._subplots.AxesSubplot at 0x267a0ae5dc8>
```



```
[36]: breast_cancer.diagnosis.value_counts()
```

```
[36]: B    357  
     M    212  
     Name: diagnosis, dtype: int64
```

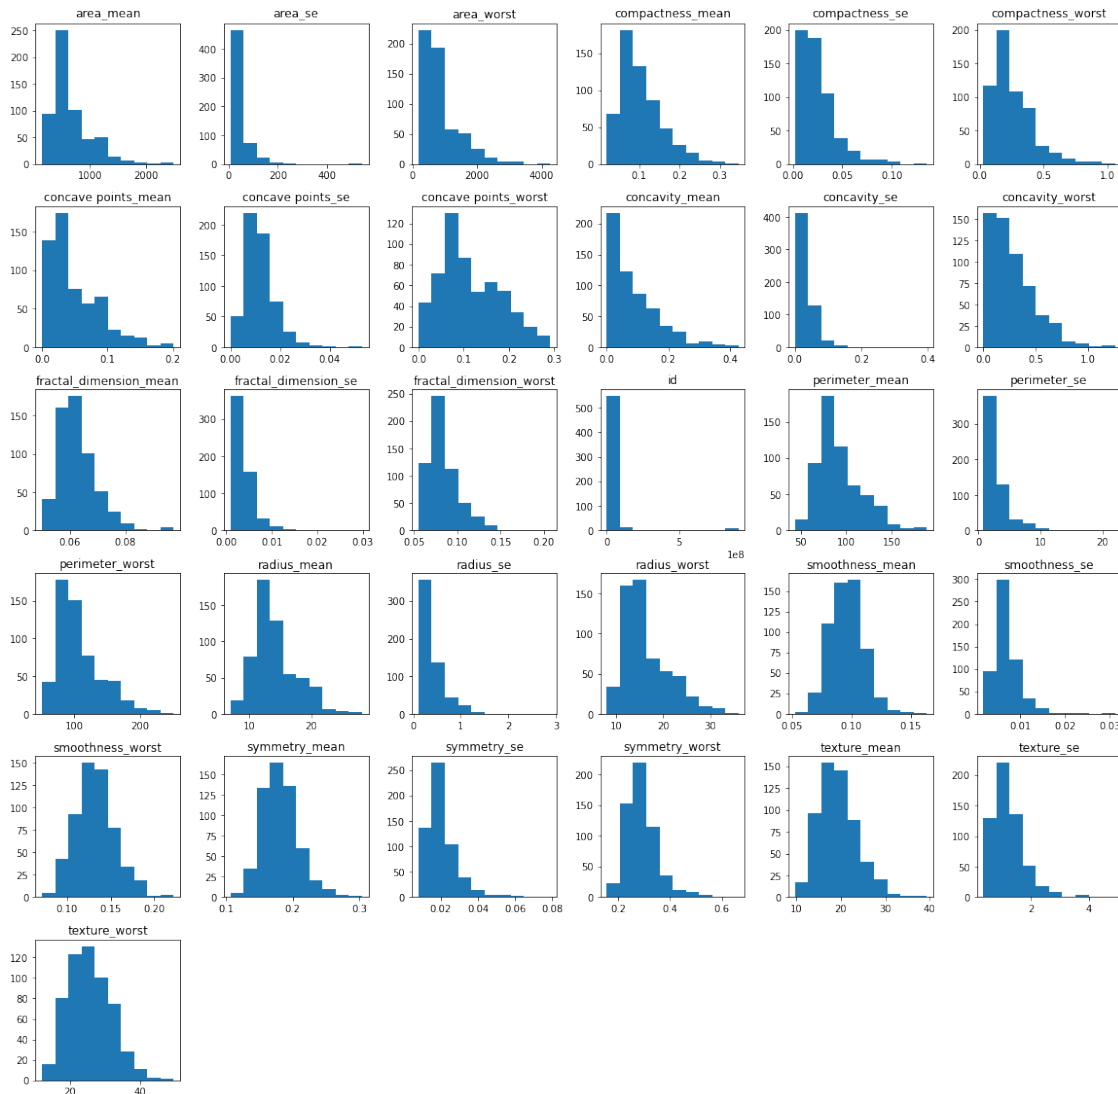
```
[37]: breast_cancer.hist(bins=10,figsize=(20,20),grid=False)
```

```
[37]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0BFC208>,  
          <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0C28148>,  
          <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0C5A5C8>,  
          <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0C920C8>,  
          <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0CBEB88>,  
          <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0CF8688>],  
        [<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0D31188>,  
         <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0D5CC88>,  
         <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0D6D508>,  
         <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0DA7148>,  
         <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0E02D08>],  
        ...])
```

```

    <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0E3CAC8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0E747C8>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0EAD548>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0EE5308>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0F1E088>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A1642DC8>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A167AB08>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0FF1308>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0ABA688>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0797048>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0AEC148>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A1064408>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0776F48>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0AD7DC8>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A169DC88>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0826B08>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A086BF88>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A089B808>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A08D4688>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000267A090E548>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A09463C8>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0982308>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A09BA248>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A09F4188>,
 <matplotlib.axes._subplots.AxesSubplot object at 0x00000267A0A2E108>]],
dtype=object)


```



```
[38]: X=breast_cancer.drop("diagnosis",axis=1)
      y=breast_cancer["diagnosis"]
```

```
[39]: features_mean= list(breast_cancer.columns[2:22])
```

```
[40]: plt.figure(figsize=(20,20))
      sns.heatmap(breast_cancer[features_mean].corr(), annot=True, square=True,
                  cmap='coolwarm')
      plt.show()
```



```
[41]: breast_cancer.drop("id", axis=1, inplace=True)
```

```
[42]: breast_cancer.head(10)
```

```
[42]:  diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0         M      17.99      10.38      122.80      1001.0
1         M      20.57      17.77      132.90      1326.0
2         M      19.69      21.25      130.00      1203.0
3         M      11.42      20.38       77.58       386.1
4         M      20.29      14.34      135.10      1297.0
5         M      12.45      15.70       82.57       477.1
6         M      18.25      19.98      119.60      1040.0
7         M      13.71      20.83       90.20       577.9
8         M      13.00      21.82       87.50       519.8
```


| | | | | | |
|---|---|-------|-------|-------|-------|
| 9 | M | 12.46 | 24.04 | 83.97 | 475.9 |
|---|---|-------|-------|-------|-------|

| | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | \ |
|---|-----------------|------------------|----------------|---------------------|---|
| 0 | 0.11840 | 0.27760 | 0.30010 | 0.14710 | |
| 1 | 0.08474 | 0.07864 | 0.08690 | 0.07017 | |
| 2 | 0.10960 | 0.15990 | 0.19740 | 0.12790 | |
| 3 | 0.14250 | 0.28390 | 0.24140 | 0.10520 | |
| 4 | 0.10030 | 0.13280 | 0.19800 | 0.10430 | |
| 5 | 0.12780 | 0.17000 | 0.15780 | 0.08089 | |
| 6 | 0.09463 | 0.10900 | 0.11270 | 0.07400 | |
| 7 | 0.11890 | 0.16450 | 0.09366 | 0.05985 | |
| 8 | 0.12730 | 0.19320 | 0.18590 | 0.09353 | |
| 9 | 0.11860 | 0.23960 | 0.22730 | 0.08543 | |

| | symmetry_mean | ... | radius_worst | texture_worst | perimeter_worst | \ |
|---|---------------|-----|--------------|---------------|-----------------|---|
| 0 | 0.2419 | ... | 25.38 | 17.33 | 184.60 | |
| 1 | 0.1812 | ... | 24.99 | 23.41 | 158.80 | |
| 2 | 0.2069 | ... | 23.57 | 25.53 | 152.50 | |
| 3 | 0.2597 | ... | 14.91 | 26.50 | 98.87 | |
| 4 | 0.1809 | ... | 22.54 | 16.67 | 152.20 | |
| 5 | 0.2087 | ... | 15.47 | 23.75 | 103.40 | |
| 6 | 0.1794 | ... | 22.88 | 27.66 | 153.20 | |
| 7 | 0.2196 | ... | 17.06 | 28.14 | 110.60 | |
| 8 | 0.2350 | ... | 15.49 | 30.73 | 106.20 | |
| 9 | 0.2030 | ... | 15.09 | 40.68 | 97.65 | |

| | area_worst | smoothness_worst | compactness_worst | concavity_worst | \ |
|---|------------|------------------|-------------------|-----------------|---|
| 0 | 2019.0 | 0.1622 | 0.6656 | 0.7119 | |
| 1 | 1956.0 | 0.1238 | 0.1866 | 0.2416 | |
| 2 | 1709.0 | 0.1444 | 0.4245 | 0.4504 | |
| 3 | 567.7 | 0.2098 | 0.8663 | 0.6869 | |
| 4 | 1575.0 | 0.1374 | 0.2050 | 0.4000 | |
| 5 | 741.6 | 0.1791 | 0.5249 | 0.5355 | |
| 6 | 1606.0 | 0.1442 | 0.2576 | 0.3784 | |
| 7 | 897.0 | 0.1654 | 0.3682 | 0.2678 | |
| 8 | 739.3 | 0.1703 | 0.5401 | 0.5390 | |
| 9 | 711.4 | 0.1853 | 1.0580 | 1.1050 | |

| | concave points_worst | symmetry_worst | fractal_dimension_worst |
|---|----------------------|----------------|-------------------------|
| 0 | 0.2654 | 0.4601 | 0.11890 |
| 1 | 0.1860 | 0.2750 | 0.08902 |
| 2 | 0.2430 | 0.3613 | 0.08758 |
| 3 | 0.2575 | 0.6638 | 0.17300 |
| 4 | 0.1625 | 0.2364 | 0.07678 |
| 5 | 0.1741 | 0.3985 | 0.12440 |
| 6 | 0.1932 | 0.3063 | 0.08368 |
| 7 | 0.1556 | 0.3196 | 0.11510 |

| | | | |
|---|--------|--------|---------|
| 8 | 0.2060 | 0.4378 | 0.10720 |
| 9 | 0.2210 | 0.4366 | 0.20750 |

[10 rows x 31 columns]

```
[43]: diagno=pd.get_dummies(breast_cancer['diagnosis'], drop_first=True)
      diagno.head()
```

```
[43]:      M
0    1
1    1
2    1
3    1
4    1
```

```
[44]: breast_cancer=pd.concat((breast_cancer,diagno),axis=1)
```

```
[45]: breast_cancer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    object
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                        569 non-null    float64
6   compactness_mean                       569 non-null    float64
7   concavity_mean                         569 non-null    float64
8   concave points_mean                    569 non-null    float64
9   symmetry_mean                          569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                              569 non-null    float64
12  texture_se                             569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                                569 non-null    float64
15  smoothness_se                          569 non-null    float64
16  compactness_se                         569 non-null    float64
17  concavity_se                           569 non-null    float64
18  concave points_se                      569 non-null    float64
19  symmetry_se                            569 non-null    float64
20  fractal_dimension_se                   569 non-null    float64
21  radius_worst                           569 non-null    float64
22  texture_worst                          569 non-null    float64
```

```

23  perimeter_worst      569 non-null    float64
24  area_worst           569 non-null    float64
25  smoothness_worst     569 non-null    float64
26  compactness_worst    569 non-null    float64
27  concavity_worst      569 non-null    float64
28  concave points_worst  569 non-null    float64
29  symmetry_worst        569 non-null    float64
30  fractal_dimension_worst 569 non-null    float64
31  M                    569 non-null    uint8

```

dtypes: float64(30), object(1), uint8(1)

memory usage: 138.5+ KB

```
[46]: breast_cancer.drop("diagnosis", axis=1, inplace=True)
```

```
[47]: breast_cancer.head()
```

```

[47]:   radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  \
0         17.99         10.38          122.80      1001.0           0.11840
1         20.57         17.77          132.90      1326.0           0.08474
2         19.69         21.25          130.00      1203.0           0.10960
3         11.42         20.38           77.58       386.1           0.14250
4         20.29         14.34          135.10      1297.0           0.10030

      compactness_mean  concavity_mean  concave points_mean  symmetry_mean  \
0           0.27760         0.3001         0.14710         0.2419
1           0.07864         0.0869         0.07017         0.1812
2           0.15990         0.1974         0.12790         0.2069
3           0.28390         0.2414         0.10520         0.2597
4           0.13280         0.1980         0.10430         0.1809

      fractal_dimension_mean  ...  texture_worst  perimeter_worst  area_worst  \
0           0.07871  ...         17.33         184.60         2019.0
1           0.05667  ...         23.41         158.80         1956.0
2           0.05999  ...         25.53         152.50         1709.0
3           0.09744  ...         26.50          98.87          567.7
4           0.05883  ...         16.67         152.20         1575.0

      smoothness_worst  compactness_worst  concavity_worst  concave points_worst  \
0           0.1622         0.6656         0.7119         0.2654
1           0.1238         0.1866         0.2416         0.1860
2           0.1444         0.4245         0.4504         0.2430
3           0.2098         0.8663         0.6869         0.2575
4           0.1374         0.2050         0.4000         0.1625

      symmetry_worst  fractal_dimension_worst  M
0           0.4601         0.11890  1
1           0.2750         0.08902  1

```

| | | | |
|---|--------|---------|---|
| 2 | 0.3613 | 0.08758 | 1 |
| 3 | 0.6638 | 0.17300 | 1 |
| 4 | 0.2364 | 0.07678 | 1 |

[5 rows x 31 columns]

```
[48]: X=breast_cancer.drop("M",axis=1)
      y=breast_cancer["M"]
```

```
[49]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,
      ↪random_state=2)
```

```
[50]: from sklearn.linear_model import LogisticRegression
      logreg = LogisticRegression(C=1e5)
      logreg.fit(X_train,y_train)
      predictions=logreg.predict(X_test)
```

C:\Users\blr0a\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:764: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
[51]: from sklearn.metrics import classification_report
      classification_report(y_test,predictions)
```

```
[51]: '          precision    recall  f1-score   support\n\n      0.92      0.94      0.93      104\n      67\n\n      accuracy          0.92      171\n      67\n\n      0.92      0.91      0.91      171\n      67\n\nweighted avg          0.92      0.92      0.92      171\n      67'
```

```
[52]: from sklearn.metrics import confusion_matrix
      confusion_matrix(y_test,predictions)
```

```
[52]: array([[98,  6],
      [ 8, 59]], dtype=int64)
```

```
[53]: from sklearn.metrics import accuracy_score
      print("train accuracy:")
```

```
print(format(logreg.score(X_train,y_train)*100.0))
print("test accuracy:")
accuracy_score(y_test,predictions)
```

train accuracy:
94.72361809045226
test accuracy:

[53]: 0.9181286549707602

```
[54]: #apply SelectKBest class to extract top 10 best features
bestfeatures = SelectKBest(score_func=chi2, k=10)
fit = bestfeatures.fit(X,y)
dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(X.columns)
#concat two dataframes for better visualization
featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores.columns = ['features','effect_score'] #naming the dataframe,
↳ columns
print(featureScores.nlargest(10,'effect_score')) #print 10 best features
```

| | features | effect_score |
|----|-----------------|---------------|
| 23 | area_worst | 112598.431564 |
| 3 | area_mean | 53991.655924 |
| 13 | area_se | 8758.504705 |
| 22 | perimeter_worst | 3665.035416 |
| 2 | perimeter_mean | 2011.102864 |
| 20 | radius_worst | 491.689157 |
| 0 | radius_mean | 266.104917 |
| 12 | perimeter_se | 250.571896 |
| 21 | texture_worst | 174.449400 |
| 1 | texture_mean | 93.897508 |

```
[82]: from sklearn.model_selection import cross_val_score
mse=cross_val_score_
↳ (logreg,X_test,y_test,scoring='neg_mean_squared_error',cv=10)
mean_mse=np.mean(mse)
print(mean_mse)
```

-0.08202614379084967

```
[58]: from sklearn.metrics import r2_score
from sklearn import metrics
from sklearn.linear_model import Ridge
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import Ridge
```

```

ridgereg = Ridge(alpha=0, normalize=True)
ridgereg.fit(X_train, y_train)
y_pred = ridgereg.predict(X_test)
print("R-Square Value",r2_score(y_test,y_pred))
print("\n")
print ("mean_absolute_error :",metrics.mean_absolute_error(y_test, y_pred))
print("\n")
print ("mean_squared_error : ",metrics.mean_squared_error(y_test, y_pred))
print("\n")
print ("root_mean_squared_error : ",np.sqrt(metrics.mean_squared_error(y_test,
↪y_pred)))

```

R-Square Value 0.7183226717497466

mean_absolute_error : 0.1963291040465723

mean_squared_error : 0.06712245214759298

root_mean_squared_error : 0.25908001109231293

2.1 Ridge Regularization

```

[83]: from sklearn.linear_model import Ridge
      from sklearn.model_selection import GridSearchCV
      ridge=Ridge()
      parameters={'alpha':
↪ [1e-15,1e-10,1e-8,1e-3,1e-2,1,5,10,20,30,35,40,45,50,55,100]}
      ridge_regression=GridSearchCV(ridge,parameters,scoring='neg_mean_squared_error',cv=5)
      ridge_regression.fit(X_train, y_train)

```

```

[83]: GridSearchCV(cv=5, estimator=Ridge(),
                param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.01, 1, 5, 10,
                20, 30, 35, 40, 45, 50, 55, 100]},
                scoring='neg_mean_squared_error')

```

```

[84]: print(ridge_regression.best_params_)
      print(ridge_regression.best_score_)

```

```

{'alpha': 0.01}
-0.05979886379489485

```

```

[85]: ridgereg = Ridge(0.01, normalize=True)
      ridgereg.fit(X_train, y_train)
      ridge_pred = ridgereg.predict(X_test)

```

```

print("R-Square Value",r2_score(y_test,ridge_pred))
print("\n")
print ("mean_absolute_error :",metrics.mean_absolute_error(y_test, y_pred))
print("\n")
print ("mean_squared_error : ",metrics.mean_squared_error(y_test, y_pred))
print("\n")
print ("root_mean_squared_error : ",np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

```

R-Square Value 0.7385366367368109


mean_absolute_error : 0.19521123955034342

mean_squared_error : 0.06517942196335559

root_mean_squared_error : 0.25530260861055765

```
[86]: sns.distplot(y_test-ridge_pred)
```

```
[86]: <matplotlib.axes._subplots.AxesSubplot at 0x267a7a8b5c8>
```



2.2 Lasso Regularization

```
[87]: from sklearn.linear_model import Lasso
      from sklearn.model_selection import GridSearchCV
      lasso=Lasso()
      parameters={'alpha':
        ↳ [1e-15,1e-10,1e-8,1e-3,1e-2,1,5,10,20,30,35,40,45,50,55,100]}
      lasso_regression=GridSearchCV(lasso,parameters,scoring='neg_mean_squared_error',cv=5)
      lasso_regression.fit(X_train, y_train)
      print(lasso_regression.best_params_)
      print(lasso_regression.best_score_)
```

```
{'alpha': 1e-08}
-0.062319742677644355
```

```
[88]: lassoreg = Lasso(1e-08, normalize=True)
      lassoreg.fit(X_train, y_train)
      lasso_pred = lassoreg.predict(X_test)
      print("R-Square Value",r2_score(y_test,lasso_pred))
      print("\n")
      print ("mean_absolute_error :",metrics.mean_absolute_error(y_test, y_pred))
      print("\n")
      print ("mean_squared_error : ",metrics.mean_squared_error(y_test, y_pred))
      print("\n")
      print ("root_mean_squared_error : ",np.sqrt(metrics.mean_squared_error(y_test,
        ↳ y_pred)))
```

```
R-Square Value 0.7264765388015957
```


```
mean_absolute_error : 0.19521123955034342
```

```
mean_squared_error : 0.06517942196335559
```

```
root_mean_squared_error : 0.25530260861055765
```

```
[89]: sns.distplot(y_test-lasso_pred)
```

```
[89]: <matplotlib.axes._subplots.AxesSubplot at 0x267a7b155c8>
```

[]:

[]: